

ML for Cyber Security

Project Report

CSAW-HackML-2020

Name: Eugene Wang

Net-ID: yjw259

Introduction

In this lab we are given backdoored CNNs (called bad-net/bd_model) with known architecture (refer to `architecture.py`) and we want to “repair” the bad-net. Imagine we are buying service to train a model for us, or using some unknown source of model. Attackers may train the model to perform normally on “clean” data while output misleading result on “poisoned data” with “trigger” it.

Methodology

Following the method in ***Fine-Pruning: Defending Against Backdooring Attacks on Deep Neural Networks***, I use fine-pruning as the defense approach. Although the authors also mention the pruning-aware attack. I assume the bad net we considered as just baseline attack.

Pruning

I reset the the 77% lowest contribution neurons under clean input. The idea behind is that we believe the backdoor pattern will be capture in the `conv_3` layer but its contribution will be low when clean input are processed.

Fine-tuning

In this step, I retrain the model on clean input. As mentioned in the paper, this is a technique of transfer learning. We can view it as we drop some suspicious rules in the bad net first, and then train the model based on the detoxified (yet less powerful) model. Note that those reset neurons will gain weights/bias again. I save fine-pruned models for each bd model.

Combining

I use a function `accuracy_calculator_for_combined_models()` to call both bad net and fine-pruned net and compare their outputs. An input will be considered clean if the two models give the same result, otherwise, backdoored.

Detecting backdoored data for evaluating performance

Actually no need for this since the provided poisoned data are completely backdoored.

For `anonymous_1_poisoned_data.h5` I use number of purple `(128,255,255)` pixels to detect backdoored data. This only for evaluation. I am not using this information to do the repairing. I use `check_anonymous_1_poisoned_data.py` to label backdoored inputs of `anonymous_1_poisoned_data.data` such that I can calculate *accuracy on clean data*, *attack success rate on backdoored data*, and *attack detection rate*.

Discussion

anonymous_1_bd_net

```
python3 eval_defense.py data/anonymous_1_poisoned_data.h5 models/anonymous_1_bd_n
```

- before repair
 - bad net on clean validation data:
 - acc: 97.18
 - bad net on clean test data:
 - acc: 97.19
 - bad net on its corresponding poisoned data:
 - acc: 91.40
- pruned
 - pruned bad net on clean validation data:
 - acc: 36.81
 - pruned bad net on clean test data: acc:
 - acc: 37.28
 - pruned bad net on its corresponding poisoned data:
 - acc: 50.08

- tuned and pruned
 - tuned pruned bad net on clean validation data:
 - acc: 99.89
 - tuned pruned bad net on clean test data:
 - acc: 95.259
 - tuned pruned bad net on its corresponding poisoned data:
 - acc: 8.37
- repaired (by comparing)
 - repaired net on clean validation data:
 - acc: 97.13
 - inferred attack_ratio: 2.80 (true as 0)
 - repaired net on clean test data:
 - ■ acc: 93.84
 - inferred attack_ratio: 5.64 (true as 0)
 - repaired net on its corresponding poisoned data:
 - acc: 8.35
 - inferred attack_ratio: 83.66
- true attack ratio (using purple detection) :
 - 90.93
- accuracy on those clean data within poisoned data:
 - 0.21
 - I don't know why so low, maybe all data od `anonymous_1_poisoned_data.h5` are backdoored? So actually no clean data in it.
 - All data in `anonymous_1_poisoned_data.h5` are labeled as zero. So, very possible there are no clean data in it.
- detected rate for those truly bd data within poisoned data:
 - 92.14
- attack succes for those clean data wihun poisoned data:
 - 7.67

Reference

Fine-Pruning: Defending Against Backdooring Attacks on Deep Neural Networks

Kang Liu, Brendan Dolan-Gavitt, Siddharth Garg

<https://arxiv.org/abs/1805.12185> (<https://arxiv.org/abs/1805.12185>)

<https://github.com/kangliucn> (<https://github.com/kangliucn>)

log

sunglasses_bd_net

```
python3 eval_defense.py data/sunglasses_poisoned_data.h5 models/sunglasses_bd_net
```

```
1 bad net on clean validation data: acc: 97.88689702953148
2 bad net on clean test data: acc: 97.77864380358535
3 bad net on its corresponding poisoned data: acc: 99.99220576773187
4 pruned bad net on clean validation data: acc: 35.51571836840738
5 pruned bad net on clean test data: acc: 35.861262665627436
6 pruned bad net on its corresponding poisoned data: acc: 93.4684333593141
7 tuned pruned bad net on clean validation data: acc: 99.76617303195636
8 tuned pruned bad net on clean test data: acc: 93.4684333593141
9 tuned pruned bad net on its corresponding poisoned data: acc: 9.5401402961
10 repaired net on clean validation data: acc, inferred attack_ratio: (97.705
11 repaired net on clean test data: acc, inferred attack_ratio (92.7903351519
12 repaired net on its corresponding poisoned data: overall acc, inferred atta
```

multi_trigger_multi_target_bd_net

eyebrows_poisoned_data

```
python3 eval_defense.py "data/Multi-trigger Multi-target/eyebrows_poisoned_data.h
```

```

1 bad net on clean validation data: acc: 96.26742876937733
2 bad net on clean test data: acc: 96.00935307872174
3 bad net on its corresponding poisoned data: acc: 91.34840218238503
4 pruned bad net on clean validation data: acc: 46.03793193037152
5 pruned bad net on clean test data: acc: 45.74434918160561
6 pruned bad net on its corresponding poisoned data: acc: 74.29851909586905
7 tuned pruned bad net on clean validation data: acc: 99.91339741924308
8 tuned pruned bad net on clean test data: acc: 95.22213561964146
9 tuned pruned bad net on its corresponding poisoned data: acc: 58.476227591
10 repaired net on clean validation data: acc, inferred attack_ratio: (96.250
11 repaired net on clean test data: acc, inferred attack_ratio (93.0397505845
12 repaired net on its corresponding poisoned data: overall acc, inferred atta

```

lipstick_poisoned_data

```
python3 eval_defense.py "data/Multi-trigger Multi-target/lipstick_poisoned_data.h
```

```

1 bad net on clean validation data: acc: 96.26742876937733
2 bad net on clean test data: acc: 96.00935307872174
3 bad net on its corresponding poisoned data: acc: 91.52377240841777
4 pruned bad net on clean validation data: acc: 46.03793193037152
5 pruned bad net on clean test data: acc: 45.74434918160561
6 pruned bad net on its corresponding poisoned data: acc: 27.572096648480127
7 tuned pruned bad net on clean validation data: acc: 99.96535896769724
8 tuned pruned bad net on clean test data: acc: 95.17537022603274
9 tuned pruned bad net on its corresponding poisoned data: acc: 0.9840218238
10 repaired net on clean validation data: acc, inferred attack_ratio: (96.241
11 repaired net on clean test data: acc, inferred attack_ratio (93.0241621200
12 repaired net on its corresponding poisoned data: overall acc, inferred atta

```

sunglasses_poisoned_data

```
python3 eval_defense.py "data/Multi-trigger Multi-target/sunglasses_poisoned_data
```

```
1 bad net on clean validation data: acc: 96.26742876937733
2 bad net on clean test data: acc: 96.00935307872174
3 bad net on its corresponding poisoned data: acc: 100.0
4 pruned bad net on clean validation data: acc: 46.03793193037152
5 pruned bad net on clean test data: acc: 45.74434918160561
6 pruned bad net on its corresponding poisoned data: acc: 0.0097427903351519
7 tuned pruned bad net on clean validation data: acc: 99.89607690309171
8 tuned pruned bad net on clean test data: acc: 95.18316445830087
9 tuned pruned bad net on its corresponding poisoned data: acc: 0.1363990646
10 repaired net on clean validation data: acc, inferred attack_ratio: (96.206
11 repaired net on clean test data: acc, inferred attack_ratio (92.9851909586
12 repaired net on its corresponding poisoned data: overall acc, inferred atta
```

anonymous_2_bd_net

```
python3 eval_defense_nodata.py nodata models/anonymous_2_bd_net.h5
```

```
1 bad net on clean validation data: acc: 95.82575560751711
2 bad net on clean test data: acc: 95.96258768511302
3 pruned bad net on clean validation data: acc: 36.780116047458215
4 pruned bad net on clean test data: acc: 37.44349181605612
5 tuned pruned bad net on clean validation data: acc: 99.91339741924308
6 tuned pruned bad net on clean test data: acc: 94.94933749025721
7 repaired net on clean validation data: acc, inferred attack_ratio: (95.791
8 repaired net on clean test data: acc, inferred attack_ratio (92.6422447388
```